

mptools

2.0.2

Generated on Mon Sep 12 2022 21:40:46 for mptools by Doxygen 1.9.5

Mon Sep 12 2022 21:40:46

1 Introduction	1
2 Installing mptools	3
2.1 Changing install location	4
3 Installing multipass	5
4 Quickstart	7
4.1 Examples	7
4.1.1 Creating nodes using naming convention	7
4.1.2 Creating nodes using the full configuration	8
5 Creating generic nodes	9
5.1 Cloud init files	9
5.2 Manually trigger creation of cloud-init files	10
5.3 Resolving location of cloud-init files	10
5.4 Examples of creating custom nodes	11
5.4.1 Setting up a Postgresql DB-server	11
6 Creating nodes using naming conventions	13
6.1 Node naming convention	13
6.1.1 MAJOR_RELEASE	13
6.1.2 CONFIG	14
6.1.3 SIZE	14
6.1.4 NODE_NUMBER	14
6.2 Node name examples	14
7 Creating nodes using make	15
7.1 Examples of using the Makefile directly	15
7.2 Makefile targets	16
8 Aliases	17
9 Tips and Tricks	19
10 ToDo	21
10.1 Planned	21
10.2 Completed	21
10.3 Will not do	21
11 File Index	23
11.1 File List	23
12 File Documentation	25
12.1 mkmpnode.sh File Reference	25
12.1.1 Detailed Description	25

12.2 mpinstall.sh File Reference	25
12.2.1 Detailed Description	26
12.3 mpn.sh File Reference	26
12.3.1 Detailed Description	26
Index	27

Chapter 1

Introduction

mptools is a utility package for **macOS** and **Linux** to help to create customized virtual machine nodes using **multipass**.

In addition it can also help facilitate an adapted installation of **multipass** using the `mpinstall` utility on **macOS**

The customization of nodes are done through `cloud-init` files. A set of template cloud-init files are provided in this package.

During installation the templates are initiated based on the current user to setup SSH keys and user name to make it easy to access the nodes from other tools (e.g. as Jenkins nodes) by simply SSH'ing into the node.

Chapter 2

Installing mptools

It is recommended to download an official release (or use a tagged version in the repo) as there is no guarantee that the latest `main` branch is ready for deployment since that by definition is work in progress.

1. Download, unpack and install the latest tar-ball `mptools-x.y.z.tar.gz`, e.g.

```
% curl -LO https://github.com/johan162/mptools/releases/download/v2.0.1/mptools-2.0.1.tar.gz
% tar xzf mptools-2.0.1.tar.gz
% cd mptools-2.0.1
% make install
```

Note: If `curl` is not installed `wget` could be used to download the package as so

```
% wget -q --show-progress
https://github.com/johan162/mptools/releases/download/v2.0.1/mptools-2.0.1.tar.gz
```

The `make install` will install the scripts under `/usr/local/bin`. The get the shell autocompletion updated either the terminal have to restarted or call `rehash` to update the shell auto-completion hash.

In addition to installing the scripts the install target will also create a hidden directory in the current users home directory at `~/.mptools`. In that directory a number of customized cloud-init files will be stored. These are customized with the current users public SSH key as well as also setting up user account with the same name as the current user in the created nodes.

This setup will then make it simple to ssh into the node for example as `% ssh 192.168.yy.xx` (where the IPv4 address is assigned to the node)

2. If `multipass` (see <https://multipass.run/>) is not installed then this is most easily done with the utility program provided

```
% mpinstall
```

Note: This requires `homebrew` to be installed and an error will be given if it is not installed.

Note

The scripts can also be run directly from the downloaded package directory (e.g. `mptools-2.0.0`). The one thing to remember is that the script files are named with the `*.sh` suffix. When the package is installed the symlink is the basename of the script without this suffix to make it slightly easier to call the script.*

2.1 Changing install location

By default, the scripts will be installed using the prefix `/usr/local` for the installation directory. This means that the package will be installed under `/usr/local/share` and the binaries will be linked in `/usr/local/bin`.

This can be changed by adjusting the `INSTALL_PREFIX` makefile variable either permanently in the `Makefile` or as an override in the call to `make`. So for example, to install into `/usr/share` and `/usr/bin`, i.e using the prefix `/usr` the following invocation would be needed:

```
% make INSTALL_PREFIX=/usr install
```

Remember, the same prefix has to be used when uninstalling the package, i.e.

```
% make INSTALL_PREFIX=/usr uninstall
```


Chapter 3

Installing multipass

The easiest way to install `multipass` is to use the provided script `mpinstall`.

```
NAME
  mpinstall.sh - Install multipass and adjust default driver
USAGE
  mpinstall.sh [-v] [-h] [-n]
SYNOPSIS
  -h          : Print help and exit
  -n          : No execution. Only display actions.
  -v          : Print version and exit
```

To install `multipass` call

```
% mpinstall
```

This will install `multipass` on a macOS (both M1 and Intel). On Intel architecture it will also replace the default *hyperkit* virtualization driver with *qemu* driver since this driver will more allow the modification of existing machine to, say, adjust the memory size.

In addition, the script will add a number of aliases to the `~/.zshenv` file to make it easier to manage and start nodes. See the section [Aliases](#) for a detailed description.

The script will automatically detect if `multipass` have already been installed, and hence can be considered idempotent.

Chapter 4

Quickstart

New nodes can now be created with either `mkmpnode` or `mpn`:

- `mkmpnode` - (*Make Multipass Node*). A generic utility to create a single node with arbitrary specifications and specified (or default) cloud-init files. Run `mkmpnode -h` for an explanation and options.
- `mpn` - (*Multipass Named Nodes*). A utility to create one or more nodes at the same time based on naming convention of nodes. The Node naming convention controls both the size of the node and the cloud-init specification used.

The relation between the scripts are shown in Figure 1.



Figure 1: Relation between scripts and the underlying *multipass*

4.1 Examples

The following basic examples show how nodes can be created once the `mptools` package have been installed. This is just a quick review without explaining all details, the details are explained in the rest of the user guide. The purpose of these examples are just to give the reader enough to determine if this package is a good fit for purpose.

4.1.1 Creating nodes using naming convention

1. `mpn ub22f101` - Create a node based on Ubuntu 22 LTS with a full development configuration (**f**) in a large (**l**) size node.
2. `mpn ub18bs01` - Create a node based on Ubuntu 18 LTS with a basic (**b**) node configuration (no development environment) in a small (**s**) node.
3. `mpn ub20ms01 ub20ms02` - Create two nodes, both based on Ubuntu 20 LTS, minimum development configuration (**m**) in a small (**s**) node.

4.1.2 Creating nodes using the full configuration

1. `mkmpnode -m 1GB mynode` - Create a node with 1GB RAM, use the default cloud-config file and name it *mynode*.
2. `mkmpnode -m 2GB -d 10GB -c pg-config.yaml db-server` - Create a node named `db-server` with Posgresql database cloud-init config with 2GB RAM and 10GB disk.

Note

In the rest of this documentation it is assumed that the package have been installed and that links to the scripts have been created in `/usr/local/bin` and that this path is included in the `$PATH` shell variable.

Chapter 5

Creating generic nodes

The one reason for this package existence is to make it easy to create nodes with cloud-init files. The main workhorse script to do so is `mkmpnode`.

```
NAME
    mkmpnode - Create multipass nodes with a specified (or default) cloud-init file
USAGE
    mkmpnode [-r RELEASE] [-c FILE] [-d SIZE] [-p CPUS] [-m SIZE] [-q] [-v] [-h] NODE_NAME
SYNOPSIS
    -r RELEASE: Valid ubuntu release [bionic focal impish jammy docker] ($ubuntuVer)
    -c FILE    : Cloud config file (${defaultCloudInit})
    -m SIZE    : Memory size, defaults (${memory})
    -d SIZE    : Disk size, defaults (${disk}GB)
    -p NUM     : Number of CPUs (${cpus})
    -M         : Mount ${HOME}/Devel inside node
    -n         : No execution. Only display actions.
    -q         : Quiet (no output to stdout)
    -v         : Print version and exit
    -h         : Print help and exit
```

Warning

Multipass does not allow a literal underscore in node names!

Most options should be self-explanatory but the `-M` deserves a comment. If the users home catalogue have a directory `~/Devel` it will be mounted automatically in the node directly under the default users (`ubuntu`) home directory.

Note

Use the `-n` flag to do a dryrun and see how the underlying call to `multipass` is made without actually executing it.

5.1 Cloud init files

As mentioned in the previous section `mkmpnode` uses cloud-init files to configure the created nodes. Cloud init files are written as human-readable YAML files.

It is out of the scope of this readme to fully describe the full syntax of cloud-init files. Instead, we refer to the official home of the cloud-init project cloud-init.io and the description of [Cloud-Init files](#).

This toolset includes a few cloud-init templates, they are all stored in the `cloud/` folder in the distributed package. As of this writing the following templates are provided:

1. `cloud/fulldev-config.in`, A full C/C++ dev environment
2. `cloud/minidev-config.in`, A minimal c/C++ dev environment
3. `cloud/mini-config.in`, A minimal node with only user and SSH keys
4. `cloud/jenkins-config.in`, A basic Jenkins node
5. `cloud/pg-config.in`, A basic Postgresql node

These template cloud-init files will be used in the installation process to create customized versions based on the current user. The generated `*.yaml` files are stored in the user home directory under `~/mptools`.

This instantiation will be done as part of the `make install` target.

5.2 Manually trigger creation of cloud-init files

For experiments, it can be handy to re-generate the cloud-init file even after they have been initially created.

Change back to the `mptools` package directory where the `Makefile` exists. Then run the default makefile target as:

```
% make
Transforming cloud/fulldev-config.in --> cloud/fulldev-config.yaml
Transforming cloud/jenkins-config.in --> cloud/jenkins-config.yaml
Transforming cloud/mini-config.in --> cloud/mini-config.yaml
Transforming cloud/minidev-config.in --> cloud/minidev-config.yaml
Transforming cloud/pg-config.in --> cloud/pg-config.yaml
Transforming cloud/sq-config.in --> cloud/sq-config.yaml
```

This will add the generated `*.yaml` files (based on the current user) together with the template files in `./cloud` directory. The created files will have the current users public SSH keys and user-name inserted.

Please note that only files where the template files has a newer modified timestamp than any existing `*.yaml` will be re-generated. To force all `*.yaml` files to be created regardless of timestamp first run `make clean`

The installed SSH keys in the nodes will make it easier for tools and "manual" access to the created nodes by simple `ssh:ing` into the nodes.

New cloud file templates can be easily added by, for example, copying an existing file to a new name and make modifications. The makefile will automatically pick up any new template files in the cloud directory and include them when instantiating cloud-init YAML files.

5.3 Resolving location of cloud-init files

Cloud init file is specified with the `-c` option. If no cloud file is specified in the call to `mkmpnode.sh` the default cloud-init file will be used, `minidev-config.yaml`. This cloud-init file installs a minimal development environment in the node.

If the cloud init file is specified with a path then that exact file and location will be used. If the file cannot be found this results in an error message.

If, however only a filename without path is specified then `mkmpnode` will search for the config file in three locations in order of priority:

1. The current working directory under the subdirectory `./cloud`
2. From the subdirectory `cloud` in the same folder where the script is located.
3. In the current users home directory under `'~/mptools'`

This priority order is used in roer to make it possible to experiment with new updated cloud-init files and have these be picked up by `mkmpnode` without having to "destroy" the original `*.yaml` files under `~/mptools`.

If the cloud-init file cannot be found in any of these places an error will be written and the script will abort.

5.4 Examples of creating custom nodes

We will start by illustrating how new nodes can be easily created with the help of the supplied `mkmpnode` script and later on we will show how the same process can be further simplified and automated by using `mpn`.

To execute these examples it is assumed that the cloud-init YAML files have been instantiated either by being installed (`make install`) or being instantiated as described above (with a call to the default `make target`).

First we are going to create a node with a custom name and more memory than default

```
% mkmpnode -m 1GB mynode
```

This will create (and start) a new node with 1GB memory named "mynode" and initialized by the default cloud-init configuration which is set to `minidev-config.yaml` in the script.

By default, the created nodes will be based on the latest Ubuntu image (i.e. Ubuntu 22 LTS a.k.a. "jammy" at the time of writing).

If we instead wanted to create a larger node, based on Ubuntu 18 LTS with a full development configuration we would instead need to call

```
% mkmpnode -r bionic -m 4GB -c fulldev-config.yaml -d 10GB mynode
```

This will create a node with 4GB RAM and a 10GB disk based on Ubuntu 18 (i.e. "bionic")

5.4.1 Setting up a Postgresql DB-server

One of the cloud-init files allow for easy setup of a postgresql server. This server needs to be created manually with the help of the `mkmpnode` script since the node naming convention (used by `mpn`) has no concept of a DB server.

The cloud init file will set up a basic postgres server with the password specified in the cloud init file. *THIS IS HIGHLY INSECURE AND IS ONLY MEANT FOR TESTING AND EXPERIMENTS!*. See table below for the actual values.

User/DB Owner	Password	DB
postgres	postgres	postgres
ubuntu	ubuntu	ubuntu_db

Table: Default roles/users created by `pg-config.in`

To create a Postgresql server (assuming the cloud yaml file have previously been instantiated) where we assume we need 2GB of RAM we could for example call

```
% mkmpnode -c pg-config.yaml -m 2GB db-server
```

The default postgresql cloud file will set up the access permission to the server so it is accessible from the outside and also create a new user "ubuntu" with default password "ubuntu" and a new DB `ubuntu_db` (for experiments). The TCP/IP access restriction is set to `"samenet"` any access must be from the same subnetwork that we are currently on (e.g. from another MP node or from the host).

Chapter 6

Creating nodes using naming conventions

To further simplify the node creation the nodes can be both created and specified in one call to `mpn` ("**M**ultipass-
Nodes") script. This is accomplished by using a specific way of naming the nodes that also instructs `mpn`
exactly how those nodes should be created.

The naming convention is described in the next section.

```
NAME
  mpn - Create multipass node by naming convention
USAGE
  mpn [-h] [-v] [-s] NODE_NAME [NODE_NAME [NODE_NAME ... ]]
SYNOPSIS
  -h      : Print help and exit
  -n      : No execution. Only display actions.
  -s      : Silent
  -v      : Print version and exit
```

To create nodes one simply specifies one or more nodes using the naming format as arguments (see next section)
for example:

```
% mpn ub18fs01 ub20m101 ub22f101
```

This will create three new nodes based on Ubuntu 18, 20 and 22 LTS images. The Ubuntu 18, and the Ubuntu 22
will both have
a full development environment in a "small" node and "large" node respectively.

The middle Ubuntu 20 based node will be a minimal development environment in a "large" node.

When creating multiple nodes the script will kick of up to four parallel node creations. This greatly reduces the total
build/creation time.

6.1 Node naming convention

```
ub<MAJOR_RELEASE><CONFIG><SIZE><NODE_NUMBER>
```

6.1.1 MAJOR_RELEASE

- 18 ("bionic")
- 20 ("focal")
- 22 ("jammy")

6.1.2 CONFIG

- `b` (=Basic node, no dev tools). Using cloud-init file: `mini-config.yaml`
- `f` (=Full dev node), Using cloud-init file: `fulldev-config.yaml`
- `m` (=Minimal dev node), Using cloud-init file: `minidev-config.yaml`

6.1.3 SIZE

- `s` (= Small=500MB RAM/5GB Disk)
- `m` (= Medium=1GB RAM/5GB Disk)
- `l` (= Large=2GB RAM/10GB Disk)
- `x` (= X-Large=4GB RAM/15GB Disk)
- `h` (= Humungous=8GB RAM/20GB Disk)

6.1.4 NODE_NUMBER

- `nn` Arbitrarily chosen two digit number to avoid name conflicts since all node names must be unique.

6.2 Node name examples

Some examples of valid names are:

- `ub20b101` - A Ubuntu 20 image, basic cloud config, large machine size
- `ub18fm01` - A Ubuntu 18 image, full development setup, medium machine size
- `ub22mx12` - A Ubuntu 22 image, minimal development setup, x-large machine size

Note

All nodes will have 2 CPUs. If more CPUs are needed then the nodes must be created with the [mkmpnode.sh](#) directly using the `-p` option.*

Chapter 7

Creating nodes using make

Note

This is only documented in order to explain some "advanced" concept in the makefile, mostly for historic reasons and for those interested in novel usage of makefiles. It is recommended to use the `mpn` script instead

The previous section showed how nodes could be manually created by giving a few parameters to the `mkmpnode` script as well as the simplified method with `mpn` using a strict naming convention of the nodes.

The makefile was the original method of creating nodes and for historic reason we finish with a short description of how this method works.

The makefile method is functionally almost identical to the method with specifically named nodes with `mpn` as described above.

This is done with the makefile target `node`.

7.1 Examples of using the Makefile directly

The makefile is used as the driver to create these named nodes. By default, the makefile has three nodes predefined which are created as so

```
% make node
```

the following three default nodes are then prepared:

- `ub18fs01` (Based on "bionic", a.k.a Ubuntu 18 LTS)
- `ub20fs01` (Based on "focal", a.k.a Ubuntu 20 LTS)
- `ub22fs01` (Based on "jammy", a.k.a Ubuntu 22 LTS)

In order to build all nodes in parallel use the usual `-j` option to make. So for example to build up to four nodes in parallel call

```
% make -j4 node
```

As their names suggest these nodes are created with a full development environment based on the cloud init template `fulldev-config.in` which installs a complete C/C++ development environment with some of the most commonly used libraries. All created machines are small.

In order to create a custom set of nodes the node names can either:

1. be supplied as overridden makefile variables (recommended) or
2. be setup by changing the makefile variable in the makefile

An example will clarify this.

Assume that we instead wanted to create two large Ubuntu 22 nodes with full development configuration and one X-Large Ubuntu 18 node with just the minimal dev environment. We can then override the `$(NODES)` makefile variable on the command line as so

```
% make NODES="ub22f111 ub22f112 ub18mx13" node
```

The makefile will "under the hood" then make the following three calls to the actual node creating script

```
./mkmpnode.sh -r jammy -c cloud/fulldev-config.yaml -m 2GB -d 10GB ub22f111
./mkmpnode.sh -r jammy -c cloud/fulldev-config.yaml -m 2GB -d 10GB ub22f112
./mkmpnode.sh -r bionic -c cloud/minidev-config.yaml -m 4GB -d 15GB ub18mx13
```

Which will create two more large "jammy" (Ubuntu 22 LTS) nodes and one x-large "bionic" (Ubuntu 18 LTS) node exactly as the node names specified.

7.2 Makefile targets

Target	Purpose
<code>all</code>	The default target that will instantiate all <code>*.yaml</code> files from the corresponding <code>*.in</code> templates.
<code>node</code>	Create the nodes specified by the <code>\$(NODES)</code> makefile variable
<code>clean</code>	Delete all generated files
<code>distclean</code>	In addition to <code>clean</code> also remove any created distribution tar-ball.
<code>dist</code>	Create a distribution tar ball
<code>install</code>	Install the package (by default <code>/usr/local</code> is used as prefix)
<code>uninstall</code>	Uninstall the package
<code>_dbg</code>	Print all Makefile variables
<code>docs</code>	Generate documentation

Chapter 8

Aliases

After the `install-script` has been run the following shell aliases will be available to save some typing.

```
alias mp="multipass"  
alias mpl="multipass list"  
alias mps="multipass shell"  
alias mpe="multipass exec"  
alias mpd="multipass delete -p"  
alias mpp="multipass purge"  
alias mpi="multipass info"  
alias mpia="multipass info --all"  
alias mpsto="multipass stop --all"  
alias mpsta="multipass start --all"  
alias mpsu="multipass suspend"  
alias mpsua="multipass suspend --all"
```

These aliases can of course also be added manually.

As an example, this will make it easy to connect to a node as so:

```
% mps ub18fs01
```

or get information on the node

```
% mpi ub22m101  
Name:                ub22m101  
State:               Running  
IPv4:                192.168.64.15  
Release:             Ubuntu 22.04.1 LTS  
Image hash:          465254c8a247 (Ubuntu 22.04 LTS)  
Load:                0.00 0.00 0.00  
Disk usage:          3.0G out of 4.7G  
Memory usage:        423.8M out of 963.3M  
Mounts:              --
```


Chapter 9

Tips and Tricks

- The logfile when creating and starting nodes are stored at `/Library/Logs/Multipass/multipassd.log` and is helpful when debugging why a node will for example not start.
- When you create many nodes the assign dynamic ip (192.168.64.xxx) can sometimes need to be reset. This is most easily done by first stopping all instances and then delete the file `/var/db/dhcpd_leases`
- Never ever use a `systemctl daemon-reload` in a cloud-init file. This will kill the SSH daemon and the multipass connection to the starting node will be lost.
- The list command (`multipass list`) have an undocumented option `--no-ipv4` to exclude the IP in the output.
- The cloud instantiation is recorded under `/var/lib/cloud` in the node. If a customized node is not working this is a good place to start troubleshooting. For example, in `/var/lib/cloud/instance/scripts/runcmd` is the run commands specified in the `RunCmd` extracted as shell commands.
- Uninstall multipass by running

```
sudo sh "/Library/Application Support/com.canonical.multipass/uninstall.sh"
or
brew uninstall --zap multipass
```
- Find available images `mp find`

Chapter 10

ToDo

10.1 Planned

- Update the mptools presentation for v2.x installation method
- Add documentation to the authors github.io pages and add a link from the README.md

10.2 Completed

- Separate the README into a proper README and then the full documentation
- Add separate script to create named nodes and not use `Makefile`
- Find directory for `*.yaml` files automatically
- Create `uninstall` target

10.3 Will not do

- Make a homebrew package with automatic installation.
This turned out to be too problematic since during install `brew` shifts to the user `brew`. This means that `make install` will pick up the wrong user to instantiate the cloud-init templates for.

Chapter 11

File Index

11.1 File List

Here is a list of all documented files with brief descriptions:

mkmpnode.sh	Mkmpnode - Setup a multipass node with options	25
mpinstall.sh	Mpinstall.sh - Install <code>multipass</code> and set up some useful aliases	25
mpn.sh	Mpn.sh - Setup nodes using naming convention	26

Chapter 12

File Documentation

12.1 mkmpnode.sh File Reference

mkmpnode - Setup a multipass node with options

Variables

- String **ubuntuVer** = jammy
- String **nodeName** =
- String **memory** = "500M"
- String **disk** = "5G"
- Integer **cpus** = 2
- String **mountDev** = 0
- String **vlist** = ("bionic" "focal" "impish" "jammy" "docker")
- Integer **noexecute** = 0
- String **cloudInit** = ""
- String **defaultCloudInit** = "minidev-config.yaml"
- Integer **quiet_flag** = 0

12.1.1 Detailed Description

mkmpnode - Setup a multipass node with options

Author

Johan Persson johan162@gmail.com

Copyright

MIT License. See LICENSE file.

12.2 mpinstall.sh File Reference

[mpinstall.sh](#) - Install `multipass` and set up some useful aliases

Variables

- Integer **quiet_flag** = 0
- Integer **noexec** = 0

12.2.1 Detailed Description

[mpinstall.sh](#) - Install `multipass` and set up some useful aliases

Author

Johan Persson johan162@gmail.com

Copyright

MIT License. See LICENSE file.

12.3 mpn.sh File Reference

[mpn.sh](#) - Setup nodes using naming convention

Variables

- String **nodecachefile** = `$(mktemp /tmp/mplist.XXXXXXXXXXXXXX)`
- String **INSTALL_PREFIX** = `"/usr/local"`
- String **INSTALL_BIN_DIR** = `"${INSTALL_PREFIX}/bin"`
- String **MKMPNODE_SCRIPT** = `"/mkmpnode.sh"`
- String **INSTALL_USERCLOUDINIT_DIR** = `"${HOME}/.mptools"`
- String **SCRIPT_DIR** = `""`
- Integer **quiet_flag** = 0
- String **nodes** = `""`
- String **nodeList** = `("")`
- Integer **noexec** = 0
- String **CLOUD_CONFIG_F** = `"fulldev-config.yaml"`
- String **CLOUD_CONFIG_B** = `"mini-config.yaml"`
- String **CLOUD_CONFIG_M** = `"minidev-config.yaml"`
- String **MACHINE_CONFIG_S** = `"500MB -d 5GB"`
- String **MACHINE_CONFIG_M** = `"1GB -d 5GB"`
- String **MACHINE_CONFIG_E** = `"3GB -d 5GB"`
- String **MACHINE_CONFIG_L** = `"2GB -d 10GB"`
- String **MACHINE_CONFIG_X** = `"4GB -d 15GB"`
- String **MACHINE_CONFIG_H** = `"8GB -d 20GB"`
- String **IMAGE_UB22** = `jammy`
- String **IMAGE_UB20** = `focal`
- String **IMAGE_UB18** = `bionic`

12.3.1 Detailed Description

[mpn.sh](#) - Setup nodes using naming convention

Author

Johan Persson johan162@gmail.com

Copyright

MIT License. See LICENSE file.

Index

mkmpnode.sh, [25](#)

mpinstall.sh, [25](#)

mpn.sh, [26](#)